

In [1]: !pip install keras

Requirement already satisfied: keras in c:\users\shadm\anaconda3\lib\site-packages (2.8.0)

In [2]: !pip install tensorflow

Requirement already satisfied: tensorflow in c:\users\shadm\anaconda3\lib\site-packages (2.8.0)

Requirement already satisfied: keras<2.9,>=2.8.0rc0 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (2.8.0)

Requirement already satisfied: astunparse>=1.6.0 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (1.6.3)

Requirement already satisfied: numpy>=1.20 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (1.21.5)

Requirement already satisfied: grpcio<2.0,>=1.24.3 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (1.42.0)

Requirement already satisfied: protobuf>=3.9.2 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (3.19.1)

Requirement already satisfied: flatbuffers>=1.12 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (2.0)

Requirement already satisfied: gast>=0.2.1 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (0.5.3)

Requirement already satisfied: absl-py>=0.4.0 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (1.0.0)

Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (0.25.0)

Requirement already satisfied: keras-preprocessing>=1.1.1 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (1.1.2)

Requirement already satisfied: termcolor>=1.1.0 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (1.1.0)

Requirement already satisfied: tf-estimator-nightly==2.8.0.dev2021122109 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (2.8.0.dev2021122109)

Requirement already satisfied: typing-extensions>=3.6.6 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (4.1.1)

Requirement already satisfied: libclang>=9.0.1 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (14.0.1)

Requirement already satisfied: wrapt>=1.11.0 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (1.12.1)

Requirement already satisfied: six>=1.12.0 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (1.16.0)

Requirement already satisfied: h5py>=2.9.0 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (3.6.0)

Requirement already satisfied: tensorboard<2.9,>=2.8 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (2.8.0)

Requirement already satisfied: google-pasta>=0.1.1 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (0.2.0)

Requirement already satisfied: setuptools in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (61.2.0)

Requirement already satisfied: opt-einsum>=2.3.2 in c:\users\shadm\anaconda3\lib\site-packages (from tensorflow) (3.3.0)

Requirement already satisfied: wheel<1.0,>=0.23.0 in c:\users\shadm\anaconda3\lib\site-packages (from astunparse>=1.6.0->tensorflow) (0.37.1)

Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in c:\users\shadm\anaconda3\lib\site-packages (from tensorboard<2.9,>=2.8->tensorflow) (0.6.1)

Requirement already satisfied: requests<3,>=2.21.0 in c:\users\shadm\anaconda3\lib\site-packages (from tensorboard<2.9,>=2.8->tensorflow) (2.27.1)

Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in c:\users\shadm\anaconda3\lib\site-packages (from tensorboard<2.9,>=2.8->tensorflow) (1.8.1)

Requirement already satisfied: werkzeug>=0.11.15 in c:\users\shadm\anaconda3\lib\site-packages (from tensorboard<2.9,>=2.8->tensorflow) (2.0.3)

Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in c:\users\shadm\anaconda3\lib\site-packages (from tensorboard<2.9,>=2.8->tensorflow) (0.4.6)

Requirement already satisfied: google-auth<3,>=1.6.3 in c:\users\shadm\anaconda3\lib\site-packages (from tensorboard<2.9,>=2.8->tensorflow) (1.33.0)

Requirement already satisfied: markdown>=2.6.8 in c:\users\shadm\anaconda3\lib\site-packages (from tensorboard<2.9,>=2.8->tensorflow) (3.3.4)

Requirement already satisfied: cachetools<5.0,>=2.0.0 in c:\users\shadm\anaconda3\lib\site-packages (from google-auth<3,>=1.6.3->tensorboard<2.9,>=2.8->tensorflow) (4.2.2)

Requirement already satisfied: rsa<5,>=3.1.4 in c:\users\shadm\anaconda3\lib\site-

packages (from google-auth<3,>=1.6.3->tensorboard<2.9,>=2.8->tensorflow) (4.7.2)
 Requirement already satisfied: pyasn1-modules>=0.2.1 in c:\users\shadm\anaconda3\lib\site-packages (from google-auth<3,>=1.6.3->tensorboard<2.9,>=2.8->tensorflow) (0.2.8)
 Requirement already satisfied: requests-oauthlib>=0.7.0 in c:\users\shadm\anaconda3\lib\site-packages (from google-auth-oauthlib<0.5,>=0.4.1->tensorboard<2.9,>=2.8->tensorflow) (1.3.1)
 Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in c:\users\shadm\anaconda3\lib\site-packages (from pyasn1-modules>=0.2.1->google-auth<3,>=1.6.3->tensorboard<2.9,>=2.8->tensorflow) (0.4.8)
 Requirement already satisfied: idna<4,>=2.5 in c:\users\shadm\anaconda3\lib\site-packages (from requests<3,>=2.21.0->tensorboard<2.9,>=2.8->tensorflow) (3.3)
 Requirement already satisfied: urllib3<1.27,>=1.21.1 in c:\users\shadm\anaconda3\lib\site-packages (from requests<3,>=2.21.0->tensorboard<2.9,>=2.8->tensorflow) (1.26.9)
 Requirement already satisfied: certifi>=2017.4.17 in c:\users\shadm\anaconda3\lib\site-packages (from requests<3,>=2.21.0->tensorboard<2.9,>=2.8->tensorflow) (2021.10.8)
 Requirement already satisfied: charset-normalizer~2.0.0 in c:\users\shadm\anaconda3\lib\site-packages (from requests<3,>=2.21.0->tensorboard<2.9,>=2.8->tensorflow) (2.0.4)
 Requirement already satisfied: oauthlib>=3.0.0 in c:\users\shadm\anaconda3\lib\site-packages (from requests-oauthlib>=0.7.0->google-auth-oauthlib<0.5,>=0.4.1->tensorboard<2.9,>=2.8->tensorflow) (3.2.0)

```
In [3]: import os

%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import os
from glob import glob
import seaborn as sns
from PIL import Image
np.random.seed(11)
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, KFold, cross_val_score, GridSearchCV
from sklearn.metrics import accuracy_score
import itertools

import keras
from keras.utils.np_utils import to_categorical
from keras.models import Sequential, Model
from keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPool2D
from keras import backend as K
from tensorflow.keras.layers import BatchNormalization
from keras.utils.np_utils import to_categorical
from tensorflow.keras.optimizers import Adam, RMSprop
from keras.preprocessing.image import ImageDataGenerator
from keras.callbacks import ReduceLROnPlateau
from keras.wrappers.scikit_learn import KerasClassifier
from tensorflow.keras.applications.resnet50 import ResNet50
from keras import backend as K
```

```
In [4]: folder_benign_train = r'C:\Users\shadm\Desktop\Project\Biomedical_Image_Processing_using_Deep_Learning\Biomedical_Image_Processing\train\benign'
folder_malignant_train = r'C:\Users\shadm\Desktop\Project\Biomedical_Image_Processing_using_Deep_Learning\Biomedical_Image_Processing\train\malignant'

folder_benign_test = r'C:\Users\shadm\Desktop\Project\Biomedical_Image_Processing_using_Deep_Learning\Biomedical_Image_Processing\test\benign'
folder_malignant_test = r'C:\Users\shadm\Desktop\Project\Biomedical_Image_Processing_using_Deep_Learning\Biomedical_Image_Processing\test\malignant'

read = lambda imname: np.asarray(Image.open(imname).convert("RGB"))
```

```

ims_benign = [read(os.path.join(folder_benign_train, filename)) for filename in os.listdir(folder_benign_train)]
X_benign = np.array(ims_benign, dtype='uint8')
ims_malignant = [read(os.path.join(folder_malignant_train, filename)) for filename in os.listdir(folder_malignant_train)]
X_malignant = np.array(ims_malignant, dtype='uint8')

ims_benign_test = [read(os.path.join(folder_benign_test, filename)) for filename in os.listdir(folder_benign_test)]
X_benign_test = np.array(ims_benign_test, dtype='uint8')
ims_malignant_test = [read(os.path.join(folder_malignant_test, filename)) for filename in os.listdir(folder_malignant_test)]
X_malignant_test = np.array(ims_malignant_test, dtype='uint8')

y_benign = np.zeros(X_benign.shape[0])
y_malignant = np.ones(X_malignant.shape[0])

y_benign_test = np.zeros(X_benign_test.shape[0])
y_malignant_test = np.ones(X_malignant_test.shape[0])

X_train = np.concatenate((X_benign, X_malignant), axis = 0)
y_train = np.concatenate((y_benign, y_malignant), axis = 0)

X_test = np.concatenate((X_benign_test, X_malignant_test), axis = 0)
y_test = np.concatenate((y_benign_test, y_malignant_test), axis = 0)

s = np.arange(X_train.shape[0])
np.random.shuffle(s)
X_train = X_train[s]
y_train = y_train[s]

s = np.arange(X_test.shape[0])
np.random.shuffle(s)
X_test = X_test[s]
y_test = y_test[s]


```

```

In [5]: w=40
        h=30
        fig=plt.figure(figsize=(12, 8))
        columns = 5
        rows = 3

        for i in range(1, columns*rows +1):
            ax = fig.add_subplot(rows, columns, i)
            if y_train[i] == 0:
                ax.title.set_text('Benign')
            else:
                ax.title.set_text('Malignant')
            plt.imshow(X_train[i], interpolation='nearest')
        plt.show()

```



```
In [6]: y_train = to_categorical(y_train, num_classes= 2)
        y_test = to_categorical(y_test, num_classes= 2)
```

```
In [7]: X_train = X_train/255.
        X_test = X_test/255.
```

```
In [8]: def build(input_shape= (224,224,3), lr = 1e-3, num_classes= 2,
        init= 'normal', activ= 'relu', optim= 'adam'):
    model = Sequential()
    model.add(Conv2D(64, kernel_size=(3, 3),padding = 'Same',input_shape=input_shape,
                    activation= activ, kernel_initializer='glorot_uniform'))
    model.add(MaxPool2D(pool_size = (2, 2)))
    model.add(Dropout(0.25))

    model.add(Conv2D(64, kernel_size=(3, 3),padding = 'Same',
                    activation =activ, kernel_initializer = 'glorot_uniform'))
    model.add(MaxPool2D(pool_size = (2, 2)))
    model.add(Dropout(0.25))

    model.add(Flatten())
    model.add(Dense(128, activation='relu', kernel_initializer=init))
    model.add(Dense(num_classes, activation='softmax'))
    model.summary()

    if optim == 'rmsprop':
        optimizer = RMSprop(lr=lr)

    else:
        optimizer = Adam(lr=lr)

    model.compile(optimizer = optimizer ,loss = "binary_crossentropy", metrics=["accuracy"])
    return model

learning_rate_reduction = ReduceLRonPlateau(monitor='val_accuracy',
                                             patience=5,
                                             verbose=1,
                                             factor=0.5,
                                             min_lr=1e-7)
```

```
In [9]: input_shape = (224,224,3)
lr = 1e-5
init = 'normal'
activ = 'relu'
optim = 'adam'
epochs = 5
batch_size = 64

model = build(lr=lr, init=init, activ=activ, optim=optim, input_shape=input_shape)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 224, 224, 64)	1792
max_pooling2d (MaxPooling2D)	(None, 112, 112, 64)	0
dropout (Dropout)	(None, 112, 112, 64)	0
conv2d_1 (Conv2D)	(None, 112, 112, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0
dropout_1 (Dropout)	(None, 56, 56, 64)	0
flatten (Flatten)	(None, 200704)	0
dense (Dense)	(None, 128)	25690240
dense_1 (Dense)	(None, 2)	258
Total params: 25,729,218		
Trainable params: 25,729,218		
Non-trainable params: 0		


C:\Users\shadm\anaconda3\lib\site-packages\keras\optimizer_v2\adam.py:105: UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.
 super(Adam, self).__init__(name, **kwargs)

```
In [10]: history = model.fit(X_train, y_train, validation_split=0.2,
                             epochs=epochs, batch_size=batch_size, verbose=1,
                             callbacks=[learning_rate_reduction])
```

```
Epoch 1/5
3/3 [=====] - 14s 3s/step - loss: 1.2237 - accuracy: 0.48
12 - val_loss: 0.8665 - val_accuracy: 0.4750 - lr: 1.0000e-05
Epoch 2/5
3/3 [=====] - 10s 3s/step - loss: 1.0302 - accuracy: 0.53
12 - val_loss: 0.8373 - val_accuracy: 0.4750 - lr: 1.0000e-05
Epoch 3/5
3/3 [=====] - 9s 3s/step - loss: 1.0347 - accuracy: 0.481
2 - val_loss: 0.8089 - val_accuracy: 0.4750 - lr: 1.0000e-05
Epoch 4/5
3/3 [=====] - 10s 3s/step - loss: 1.0882 - accuracy: 0.42
50 - val_loss: 0.7539 - val_accuracy: 0.4750 - lr: 1.0000e-05
Epoch 5/5
3/3 [=====] - 9s 3s/step - loss: 1.1247 - accuracy: 0.512
5 - val_loss: 0.7295 - val_accuracy: 0.4750 - lr: 1.0000e-05
```

```
In [11]: print(history.history.keys())
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy', 'lr'])
```



```
In [12]: K.clear_session()
del model
del history
```

```
In [13]: kfold = KFold(n_splits=3, shuffle=True, random_state=11)

cvscores = []
for train, test in kfold.split(X_train, y_train):
    model = build(lr=lr,
                  init= init,
                  activ= activ,
```

```
        optim=optim,
        input_shape= input_shape)

model.fit(X_train[train], y_train[train], epochs=epochs, batch_size=batch_size)
scores = model.evaluate(X_train[test], y_train[test], verbose=0)
print("%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
cvscores.append(scores[1] * 100)
K.clear_session()
del model

print("%.2f%% (+/- %.2f%%)" % (np.mean(cvscores), np.std(cvscores)))
```


Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 224, 224, 64)	1792
max_pooling2d (MaxPooling2D)	(None, 112, 112, 64)	0
dropout (Dropout)	(None, 112, 112, 64)	0
conv2d_1 (Conv2D)	(None, 112, 112, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0
dropout_1 (Dropout)	(None, 56, 56, 64)	0
flatten (Flatten)	(None, 200704)	0
dense (Dense)	(None, 128)	25690240
dense_1 (Dense)	(None, 2)	258

=====
Total params: 25,729,218
Trainable params: 25,729,218
Non-trainable params: 0

Epoch 1/5

3/3 [=====] - 10s 2s/step - loss: 1.3326 - accuracy: 0.5714

Epoch 2/5

3/3 [=====] - 8s 2s/step - loss: 1.2247 - accuracy: 0.5188

Epoch 3/5

3/3 [=====] - 8s 2s/step - loss: 1.2222 - accuracy: 0.4887

Epoch 4/5

3/3 [=====] - 9s 2s/step - loss: 1.2725 - accuracy: 0.5188

Epoch 5/5

3/3 [=====] - 8s 2s/step - loss: 1.0865 - accuracy: 0.5113

accuracy: 52.24%

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 224, 224, 64)	1792
max_pooling2d (MaxPooling2D)	(None, 112, 112, 64)	0
dropout (Dropout)	(None, 112, 112, 64)	0
conv2d_1 (Conv2D)	(None, 112, 112, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0
dropout_1 (Dropout)	(None, 56, 56, 64)	0
flatten (Flatten)	(None, 200704)	0

dense (Dense) (None, 128) 25690240

dense_1 (Dense) (None, 2) 258

=====

Total params: 25,729,218
Trainable params: 25,729,218
Non-trainable params: 0

Epoch 1/5

3/3 [=====] - 10s 2s/step - loss: 0.9418 - accuracy: 0.4887

Epoch 2/5

3/3 [=====] - 8s 2s/step - loss: 0.9604 - accuracy: 0.5940

Epoch 3/5

3/3 [=====] - 8s 2s/step - loss: 0.9396 - accuracy: 0.5338

Epoch 4/5

3/3 [=====] - 8s 2s/step - loss: 0.9063 - accuracy: 0.5188

Epoch 5/5

3/3 [=====] - 8s 2s/step - loss: 0.9199 - accuracy: 0.5338

accuracy: 52.24%

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 224, 224, 64)	1792
max_pooling2d (MaxPooling2D)	(None, 112, 112, 64)	0
dropout (Dropout)	(None, 112, 112, 64)	0
conv2d_1 (Conv2D)	(None, 112, 112, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0
dropout_1 (Dropout)	(None, 56, 56, 64)	0
flatten (Flatten)	(None, 200704)	0
dense (Dense)	(None, 128)	25690240
dense_1 (Dense)	(None, 2)	258

=====

Total params: 25,729,218
Trainable params: 25,729,218
Non-trainable params: 0

Epoch 1/5

3/3 [=====] - 10s 2s/step - loss: 2.1285 - accuracy: 0.5597

Epoch 2/5

3/3 [=====] - 9s 2s/step - loss: 1.3042 - accuracy: 0.5149

Epoch 3/5

3/3 [=====] - 8s 2s/step - loss: 1.4246 - accuracy: 0.5821

```
Epoch 4/5
3/3 [=====] - 9s 2s/step - loss: 1.5070 - accuracy: 0.477
6
Epoch 5/5
3/3 [=====] - 8s 2s/step - loss: 1.1514 - accuracy: 0.619
4
accuracy: 48.48%
50.99% (+/- 1.77%)
```

```
In [14]: model = build(lr=lr,
                        init=init,
                        activ=activ,
                        optim=optim,
                        input_shape= input_shape)

model.fit(X_train, y_train,
          epochs=epochs, batch_size= batch_size, verbose=1,
          callbacks=[learning_rate_reduction]
        )

y_predict = model.predict_classes(X_test)

print(accuracy_score(np.argmax(y_test, axis=1),y_predict))
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 224, 224, 64)	1792
max_pooling2d (MaxPooling2D)	(None, 112, 112, 64)	0
dropout (Dropout)	(None, 112, 112, 64)	0
conv2d_1 (Conv2D)	(None, 112, 112, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0
dropout_1 (Dropout)	(None, 56, 56, 64)	0
flatten (Flatten)	(None, 200704)	0
dense (Dense)	(None, 128)	25690240
dense_1 (Dense)	(None, 2)	258

Total params: 25,729,218
 Trainable params: 25,729,218
 Non-trainable params: 0

Epoch 1/5

4/4 [=====] - ETA: 0s - loss: 1.1890 - accuracy: 0.4650
 WARNING:tensorflow:Learning rate reduction is conditioned on metric `val_accuracy` which is not available. Available metrics are: loss,accuracy,lr
 4/4 [=====] - 13s 3s/step - loss: 1.1890 - accuracy: 0.4650 - lr: 1.0000e-05

Epoch 2/5

4/4 [=====] - ETA: 0s - loss: 1.1305 - accuracy: 0.4850
 WARNING:tensorflow:Learning rate reduction is conditioned on metric `val_accuracy` which is not available. Available metrics are: loss,accuracy,lr
 4/4 [=====] - 12s 3s/step - loss: 1.1305 - accuracy: 0.4850 - lr: 1.0000e-05

Epoch 3/5

4/4 [=====] - ETA: 0s - loss: 1.0075 - accuracy: 0.5100
 WARNING:tensorflow:Learning rate reduction is conditioned on metric `val_accuracy` which is not available. Available metrics are: loss,accuracy,lr
 4/4 [=====] - 12s 3s/step - loss: 1.0075 - accuracy: 0.5100 - lr: 1.0000e-05

Epoch 4/5

4/4 [=====] - ETA: 0s - loss: 0.9106 - accuracy: 0.6250
 WARNING:tensorflow:Learning rate reduction is conditioned on metric `val_accuracy` which is not available. Available metrics are: loss,accuracy,lr
 4/4 [=====] - 11s 3s/step - loss: 0.9106 - accuracy: 0.6250 - lr: 1.0000e-05

Epoch 5/5

4/4 [=====] - ETA: 0s - loss: 0.8984 - accuracy: 0.6400
 WARNING:tensorflow:Learning rate reduction is conditioned on metric `val_accuracy` which is not available. Available metrics are: loss,accuracy,lr
 4/4 [=====] - 11s 3s/step - loss: 0.8984 - accuracy: 0.6400 - lr: 1.0000e-05

```

-----
AttributeError                                Traceback (most recent call last)
Input In [14], in <cell line: 12>()
      1 model = build(lr=lr,
      2               init= init,
      3               activ= activ,
      4               optim=optim,
      5               input_shape= input_shape)
      7 model.fit(X_train, y_train,
      8             epochs=epochs, batch_size= batch_size, verbose=1,
      9             callbacks=[learning_rate_reduction]
     10             )
--> 12 y_predict = model.predict_classes(X_test)
     14 print(accuracy_score(np.argmax(y_test, axis=1),y_predict))

AttributeError: 'Sequential' object has no attribute 'predict_classes'

```

```

In [15]: model_json = model.to_json()

with open("model.json", "w") as json_file:
    json_file.write(model_json)

model.save_weights("model.h5")
print("Saved model to disk")

del model
K.clear_session()

```

Saved model to disk

```

In [16]: input_shape = (224,224,3)
lr = 1e-5
epochs = 3
batch_size = 64

model = ResNet50(include_top=True,
                  weights= None,
                  input_tensor=None,
                  input_shape=input_shape,
                  pooling='avg',
                  classes=2)

model.compile(optimizer = Adam(lr) ,
              loss = "binary_crossentropy",
              metrics=["accuracy"])

history = model.fit(X_train, y_train, validation_split=0.2,
                    epochs= epochs, batch_size= batch_size, verbose=1,
                    callbacks=[learning_rate_reduction]
                    )

print(history.history.keys())
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')

```

```
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

Epoch 1/3



3/3 [=====] - ETA: 0s - loss: 0.8391 - accuracy: 0.5063 W
ARNING:tensorflow:5 out of the last 15 calls to <function Model.make_test_function.
<locals>.test_function at 0x000001068A9A2E50> triggered tf.function retracing. T
racing is expensive and the excessive number of tracings could be due to (1) creat
ing @tf.function repeatedly in a loop, (2) passing tensors with different shapes,
(3) passing Python objects instead of tensors. For (1), please define your @tf.fun
ction outside of the loop. For (2), @tf.function has experimental_relax_shapes=Tru
e option that relaxes argument shapes that can avoid unnecessary retracing. For
(3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing
and https://www.tensorflow.org/api_docs/python/tf/function for more details.
3/3 [=====] - 65s 16s/step - loss: 0.8391 - accuracy: 0.5
063 - val_loss: 0.6927 - val_accuracy: 0.5250 - lr: 1.0000e-05

Epoch 2/3

3/3 [=====] - 51s 16s/step - loss: 0.7631 - accuracy: 0.5
063 - val_loss: 0.6927 - val_accuracy: 0.5250 - lr: 1.0000e-05

Epoch 3/3

3/3 [=====] - 53s 16s/step - loss: 0.7197 - accuracy: 0.5
063 - val_loss: 0.6934 - val_accuracy: 0.5250 - lr: 1.0000e-05
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy', 'lr'])



```
In [17]: y_pred = model.predict(X_test)
print(accuracy_score(np.argmax(y_test, axis=1), np.argmax(y_pred, axis=1)))

resnet50_json = model.to_json()
```

```
with open("resnet50.json", "w") as json_file:  
    json_file.write(resnet50_json)
```

```
model.save_weights("resnet50.h5")  
print("Saved model to disk")
```

0.485

Saved model to disk

```
In [18]: del model  
K.clear_session()
```

```
In [19]: from keras.applications.vgg16 import VGG16  
from keras.preprocessing import image  
from keras.applications.vgg16 import preprocess_input  
import numpy as np  
  
input_shape = (224,224,3)  
lr = 1e-5  
epochs = 3  
batch_size = 64  
  
model = VGG16(include_top=True,  
              weights=None,  
              input_tensor=None,  
              input_shape=input_shape,  
              pooling='avg',  
              classes=2)  
  
model.compile(optimizer = Adam(lr) ,  
              loss = "binary_crossentropy",  
              metrics=["accuracy"])  
  
history = model.fit(X_train, y_train, validation_split=0.2,  
                   epochs= epochs, batch_size= batch_size, verbose=1,  
                   callbacks=[learning_rate_reduction]  
                   )  
  
print(history.history.keys())  
plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])  
plt.title('model accuracy')  
plt.ylabel('accuracy')  
plt.xlabel('epoch')  
plt.legend(['train', 'test'], loc='upper left')  
plt.show()  
plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])  
plt.title('model loss')  
plt.ylabel('loss')  
plt.xlabel('epoch')  
plt.legend(['train', 'test'], loc='upper left')  
plt.show()
```

Epoch 1/3



3/3 [=====] - ETA: 0s - loss: 0.6931 - accuracy: 0.4625 W
 ARNING:tensorflow:5 out of the last 13 calls to <function Model.make_test_function.<locals>.test_function at 0x0000010672A6C550> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has experimental_relax_shapes=True option that relaxes argument shapes that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.
 3/3 [=====] - 115s 34s/step - loss: 0.6931 - accuracy: 0.4625 - val_loss: 0.6929 - val_accuracy: 0.4750 - lr: 1.0000e-05

Epoch 2/3

3/3 [=====] - 132s 40s/step - loss: 0.6929 - accuracy: 0.5063 - val_loss: 0.6925 - val_accuracy: 0.6750 - lr: 1.0000e-05

Epoch 3/3

3/3 [=====] - 111s 34s/step - loss: 0.6926 - accuracy: 0.7812 - val_loss: 0.6920 - val_accuracy: 0.8750 - lr: 1.0000e-05
 dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy', 'lr'])



```
In [20]: y_pred = model.predict(X_test)
print(accuracy_score(np.argmax(y_test, axis=1), np.argmax(y_pred, axis=1)))

vgg16_json = model.to_json()

with open("vgg16.json", "w") as json_file:
    json_file.write(vgg16_json)
```



```
model.save_weights("vgg16.h5")
print("Saved model to disk")
```

```
del model
K.clear_session()
```

0.725

Saved model to disk

In []:

In []:

In [21]: `from keras.applications.xception import Xception`

```
input_shape = (224,224,3)
lr = 1e-5
epochs = 3
batch_size = 64
```

```
model = Xception(include_top=True,
                  weights=None,
                  input_tensor=None,
                  input_shape=input_shape,
                  pooling='avg',
                  classes=2)
```

```
model.compile(optimizer = Adam(lr) ,
              loss = "binary_crossentropy",
              metrics=["accuracy"])

history = model.fit(X_train, y_train, validation_split=0.2,
                   epochs= epochs, batch_size= batch_size, verbose=1,
                   callbacks=[learning_rate_reduction]
                   )
```

```
print(history.history.keys())
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

Epoch 1/3

3/3 [=====] - 167s 44s/step - loss: 0.6884 - accuracy: 0.5063 - val_loss: 0.6931 - val_accuracy: 0.5250 - lr: 1.0000e-05


Epoch 2/3

3/3 [=====] - 152s 45s/step - loss: 0.6379 - accuracy: 0.5813 - val_loss: 0.6931 - val_accuracy: 0.5250 - lr: 1.0000e-05

Epoch 3/3

3/3 [=====] - 172s 47s/step - loss: 0.5993 - accuracy: 0.7500 - val_loss: 0.6931 - val_accuracy: 0.5250 - lr: 1.0000e-05

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy', 'lr'])
```



```
In [22]: y_pred = model.predict(X_test)
print(accuracy_score(np.argmax(y_test, axis=1), np.argmax(y_pred, axis=1)))

xception_json = model.to_json()

with open("xception.json", "w") as json_file:
    json_file.write(xception_json)

model.save_weights("xception.h5")
print("Saved model to disk")

del model
K.clear_session()

0.5
Saved model to disk
```

```
In [ ]:
```